



GILLINGS SCHOOL OF  
GLOBAL PUBLIC HEALTH



# Introduction to Q-learning

Yang Ruan, Yuanyuan Yan, Yajie He

04/12/2024

# Content

1: Basic setting in reinforcement learning

2: *Intro to Q-learning*

3: *Case study: Q-learning in PM*

# Basic Concept

- **State:** Describes the agent's status with respect to the environment.  
**State Space:** The set of all the states, denote as  $\mathcal{S} = \{s_1, \dots, s_9\}$
- **Action:** For each state, the possible action the agent can take. e.g. moving upward, moving rightward.  
**Action Space:** The set of all actions, denoted as  $\mathcal{A} = \{a_1, \dots, a_5\}$
- **Policy:** A policy tells the agent which actions to take at every state, policies can be described by conditional probabilities.


$$\pi(a_1 \mid s_1) = 0$$

$$\pi(a_2 \mid s_1) = 1$$

# Reward and Trajectory

- **Reward**  $R_t$ : After executing an action at a state, the agent obtains a reward, denoted as  $r$ , as feedback from the environment.
- **Trajectory**: A state-action-reward chain.

$$\{S_1 = s_1, A_1 = a_2, R_1 = 0, S_2 = s_2, A_2 = a_2, R_2 = 0, S_3 = s_5, \dots\}$$




$$s_1 \xrightarrow[r=0]{a_2} s_2 \xrightarrow[r=0]{a_2} s_5 \xrightarrow[r=0]{a_3} s_8 \xrightarrow[r=1]{a_2} s_9.$$

# Return and discount rate

- **Return**  $G_t$ : The sum of all the rewards collected along the trajectory.

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = r_{t+1} + \gamma G_{t+1}$$

- Here we introduce a discount rate  $\gamma \in [0, 1)$ :
  - (1) If  $\gamma$  is close to 0, the value of the discounted return is dominated by the rewards obtained in the **near future**.
  - (2) If  $\gamma$  is close to 1, the value of the discounted return is dominated by the rewards obtained in the **far future**.




# Markov Decision Process (MDP)

- We define the state transition probability as:

$$p(s' | s, a) \doteq \Pr \{S_t = s' \mid S_{t-1} = s, A_{t-1} = a\}$$

- A Markov decision process is:



# State-Value function and Action-Value function

- **State-Value function**

state value is defined as

$$V(s) \doteq \mathbb{E}[G_t \mid S_t = s]$$

which is the expected return one would get if starting from state  $s$

- **Action-Value function**

The action value of a state-action pair  $(s, a)$  is defined as

$$q(s, a) \doteq \mathbb{E}[G_t \mid S_t = s, A_t = a]$$

which is the expected return one would get if starting from state  $s$  and take action  $a$



# Bellman Optimality Equation (BOE)

- The Bellman equation is obtained by conditional expectation:




$$\begin{aligned}Q_{\pi}(s, a) &= \mathbb{E}[G_t \mid s_t = s, a_t = a] \\&= \mathbb{E}[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \mid s_t = s, a_t = a] \\&= \mathbb{E}[r_{t+1} \mid s_t = s, a_t = a] + \gamma \mathbb{E}[r_{t+2} + \gamma r_{t+3} + \gamma^2 r_{t+4} + \dots \mid s_t = s, a_t = a] \\&= R(s, a) + \gamma \mathbb{E}[G_{t+1} \mid s_t = s, a_t = a] \\&= R(s, a) + \gamma \mathbb{E}[V(s_{t+1}) \mid s_t = s, a_t = a] \\&= R(s, a) + \gamma \sum_{s' \in S} p(s' \mid s, a) V_{\pi}(s')\end{aligned}$$

- Note that  $V_{\pi}(s) = \mathbb{E}_{\pi}[G_t \mid s_t = s] = \sum_{a \in A} \pi(a \mid s) Q_{\pi}(s, a)$ , we can have two Bellman optimality equations by plug-in either  $Q$  or  $V$ :

$$\begin{aligned}V_{\pi}(s) &= \sum_{a \in A} \pi(a \mid s) \left( R(s, a) + \gamma \sum_{s' \in S} p(s' \mid s, a) V_{\pi}(s') \right) \\Q_{\pi}(s, a) &= R(s, a) + \gamma \sum_{s' \in S} p(s' \mid s, a) \sum_{a' \in A} \pi(a' \mid s') Q_{\pi}(s', a')\end{aligned}$$



# Principle of optimality



Based on the principle,

$V_\pi(s)$  and  $Q_\pi(s, a)$  achieve the optimality at the same time.

$V_\pi(s)$  and  $\pi_k$  achieve the optimality at the same time.

# Policy Iteration

- An important algorithm for optimizing policy would be **policy iteration**, which is closely related to Bellman optimality equation as well as Markov decision process.
- Policy evaluation step: this step evaluates a given policy by calculating the corresponding state value. That is to solve the following Bellman equation:

$$V_{\pi_k}(s) = \sum_{a \in A} \pi(a | s) \left( R(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V_{\pi_k}(s') \right)$$

where  $\pi_k$  is the policy obtained in the  $k$ -th iteration and  $V_{\pi_k}$  is the state value to be calculated.

- Policy improvement step: this step aims to improve the policy. The updated policy  $\pi_{k+1}$  based on policy evaluation step can be obtained by

$$a_k^*(s) = \arg \max_a q_{\pi_k}(s, a)$$
$$\pi_{k+1}(a | s) = \begin{cases} 1, & a = a_k^*(s) \\ 0, & a \neq a_k^*(s) \end{cases}$$

# Policy Iteration

- One can show that the state value sequence  $\{V_{\pi_k}\}_{k=0}^{\infty}$  generated by the policy iteration algorithm converges to the optimal state value  $V^*$ . As a result, the policy sequence  $\{\pi_k\}_{k=0}^{\infty}$  converges to an optimal policy  $\pi^*$ .

---

## Algorithm Policy Iteration Algorithm

---

- 1: **Initialization:** System model,  $p(r \mid s, a)$  and  $p(s' \mid s, a)$  for all  $(s, a)$ , initial policy  $\pi_0$
- 2: **Goal:** Search for the optimal state value and an optimal policy.
- 3: **while**  $V_{\pi_k}$  has not converged, for the  $k$ -th iteration **do**
- 4:     **Policy Evaluation:**
- 5:     Initialization: an arbitrary  $V_{\pi_k}^{(0)}$
- 6:     **while**  $V_{\pi_k}^{(j)}$  has not converged, for the  $j$ -th iteration **do**
- 7:         **for** each state  $s \in \mathcal{S}$  **do**
- 8:              $V_{\pi_k}^{(j+1)}(s) \leftarrow \sum_a \pi_k(a \mid s) \left[ \sum_r p(r \mid s, a) r + \gamma \sum_{s'} p(s' \mid s, a) V_{\pi_k}^{(j)}(s') \right]$
- 9:         **Policy Improvement:**
- 10:         **for** each state  $s \in \mathcal{S}$  **do**
- 11:             **for** each action  $a \in \mathcal{A}$  **do**
- 12:                  $q_{\pi_k}(s, a) \leftarrow \sum_r p(r \mid s, a) r + \gamma \sum_{s'} p(s' \mid s, a) v_{\pi_k}(s')$
- 13:                  $a_k^*(s) \leftarrow \arg \max_a q_{\pi_k}(s, a)$
- 14:                  $\pi_{k+1}(a \mid s) \leftarrow 1$  if  $a = a_k^*$  and  $\pi_{k+1}(a \mid s) \leftarrow 0$  otherwise

# Content

1: *Basic setting in reinforcement learning*

2: **Intro to Q-learning**

3: *Case study: Q-learning in PM*

# Motivating Example

From table to function

- As for discussion in previous slides, state and action values are represented by tables.


State	$s_1$	$s_2$	$\dots$	$s_n$
True value	$v_\pi(s_1)$	$v_\pi(s_2)$	$\dots$	$v_\pi(s_n)$
Estimated value	$\hat{v}(s_1)$	$\hat{v}(s_2)$	$\dots$	$\hat{v}(s_n)$

	$a_1$	$a_2$	$a_3$	$a_4$	$a_5$
$s_1$	$q_\pi(s_1, a_1)$	$q_\pi(s_1, a_2)$	$q_\pi(s_1, a_3)$	$q_\pi(s_1, a_4)$	$q_\pi(s_1, a_5)$
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
$s_9$	$q_\pi(s_9, a_1)$	$q_\pi(s_9, a_2)$	$q_\pi(s_9, a_3)$	$q_\pi(s_9, a_4)$	$q_\pi(s_9, a_5)$

- Advantage: intuitive and easy to analyze
- Disadvantage: difficult to handle large or continuous state or action spaces.
  - storage
  - generalization ability

# Motivating Example

- For example, we can use a simple straight line to fit the dots. Suppose the equation of the straight line is



$$\hat{v}(s, w) = as + b = \underbrace{[s, 1]}_{\phi^T(s)} \underbrace{\begin{bmatrix} a \\ b \end{bmatrix}}_w = \phi^T(s)w$$

where  $w$  is the parameter vector;  $\phi(s)$  the feature vector of  $s$ ;  $\hat{v}(s, w)$  is linear in  $w$ .


**Benefit** storage: we do not need to store  $|\mathcal{S}|$  state values. We only need to store a lower-dimensional  $w$ .

# Motivating Example

- How do we update the state-value function?



(a) Tabular method: when  $\hat{v}(s_3)$  is updated, the other values remain the same.



(b) Function approximation method: when we update  $\hat{v}(s_3)$  by changing  $w$ , the values of the neighboring states are also changed.

$$\hat{v}(s, w) = \phi^T(s)w$$

**Benefit** generalization ability: When we update  $\hat{v}(s)$  by changing  $w$ , the values of the neighboring states are also changed.

# Stochastic gradient for objective function

- The objective function is

$$J(\mathbf{w}) = \mathbb{E} \left[ (v_\pi(S) - \hat{v}(S, \mathbf{w}))^2 \right].$$

- Our goal is to find the best  $\mathbf{w}$  that can minimize  $J(\mathbf{w})$ .
- To minimize the objective function  $J(\mathbf{w})$ , we can use the gradient-descent algorithm:

$$\begin{aligned} \mathbf{w}_{k+1} &= \mathbf{w}_k - \alpha_k \nabla_{\mathbf{w}} J(\mathbf{w}_k) \\ \longrightarrow \mathbf{w}_{k+1} &= \mathbf{w}_k + 2\alpha_k \mathbb{E}_S [(v_\pi(S) - \hat{v}(S, \mathbf{w}_k)) \nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w}_k)] \end{aligned}$$

- However, the true gradient above involves the calculation of an expectation. We can use the stochastic gradient to replace the true gradient:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha_t (v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)) \nabla_{\mathbf{w}} \hat{v}(s_t, \mathbf{w}_t) \quad (1)$$

- Notably, (2) is not implementable because it requires the true state value  $v$ , which is unknown and must be estimated.



# Temporal-difference with function approximation

- Note that  $v(s_t) = r_{t+1} + \gamma v(s_{t+1})$ , we replace  $v(s_{t+1})$  with its estimate and the algorithm becomes

$$w_{t+1} = w_t + \alpha_t [r_{t+1} + \gamma \hat{v}(s_{t+1}, w_t) - \hat{v}(s_t, w_t)] \nabla_w \hat{v}(s_t, w_t)$$

- This is the algorithm of TD learning with function approximation. An algorithm can be summarized as:

---

## Algorithm TD learning of state values with function approximation

---

- 1: **Initialization:** A function  $\hat{v}(s, w)$  that is differentiable in  $w$ . Initial parameter  $w_0$ .
  - 2: **Goal:** Learn the true state values of a given policy  $\pi$ .
  - 3: **for** each trajectory  $\{(s_t, r_{t+1}, s_{t+1})\}_t$  generated by  $\pi$  **do**
  - 4:     **for** each sample  $(s_t, r_{t+1}, s_{t+1})$  **do**
  - 5:         In the general case,  $w_{t+1} = w_t + \alpha_t [r_{t+1} + \gamma \hat{v}(s_{t+1}, w_t) - \hat{v}(s_t, w_t)] \nabla_w \hat{v}(s_t, w_t)$
  - 6:         In the linear case,  $w_{t+1} = w_t + \alpha_t [r_{t+1} + \gamma \phi^T(s_{t+1}) w_t - \phi^T(s_t) w_t] \phi(s_t)$
  - 7:     **end for**
  - 8: **end for**
-

# Sarsa with function approximation

- Note that by principle of optimality, we can intuitively replace the estimate of  $v$  with  $q$ . The Sarsa algorithm with value function approximation is

$$w_{t+1} = w_t + \alpha_t [r_{t+1} + \gamma \hat{q}(s_{t+1}, a_{t+1}, w_t) - \hat{q}(s_t, a_t, w_t)] \nabla_w \hat{q}(s_t, a_t, w_t)$$

---

## Algorithm Sarsa with function approximation

---

- 1: **Aim:** Search a policy that can lead the agent to the target from an initial state-action pair  $(s_0, a_0)$ .
  - 2: **for** each trajectory **do**
  - 3:     **if** the current  $s_t$  is not the target state **then**
  - 4:         Take action  $a_t$  following  $\pi_t(s_t)$ , generate  $r_{t+1}, s_{t+1}$ , and then take action  $a_{t+1}$  following  $\pi_t(s_{t+1})$
  - 5:         **Value update (parameter update):**  
$$w_{t+1} = w_t + \alpha_t [r_{t+1} + \gamma \hat{q}(s_{t+1}, a_{t+1}, w_t) - \hat{q}(s_t, a_t, w_t)] \nabla_w \hat{q}(s_t, a_t, w_t)$$
  - 6:         **Policy update:**  
$$\pi_{t+1}(a | s_t) = \begin{cases} 1 - \frac{\epsilon}{|\mathcal{A}(s)|} (|\mathcal{A}(s)| - 1) & \text{if } a = \arg \max_{a \in \mathcal{A}(s_t)} \hat{q}(s_t, a, w_{t+1}) \\ \frac{\epsilon}{|\mathcal{A}(s)|} & \text{otherwise} \end{cases}$$
  - 7:     **end if**
  - 8: **end for**
-

# Q-learning with function approximation

- The Q-learning algorithm with value function approximation is

$$w_{t+1} = w_t + \alpha_t \left[ r_{t+1} + \gamma \max_{a \in \mathcal{A}(s_{t+1})} \hat{q}(s_{t+1}, a, w_t) - \hat{q}(s_t, a_t, w_t) \right] \nabla_w \hat{q}(s_t, a_t, w_t)$$

---

## Algorithm Q-learning with function approximation (on-policy version)

---

- 1: **Initialization:** Initial parameter vector  $w_0$ . Initial policy  $\pi_0$ . Small  $\varepsilon > 0$ .
- 2: **Aim:** Search a good policy that can lead the agent to the target from an initial state-action pair  $(s_0, a_0)$ .
- 3: **for** each trajectory **do**
- 4:     **if** the current  $s_t$  is not the target state **then**
- 5:         Take action  $a_t$  following  $\pi_t(s_t)$ , and generate  $r_{t+1}, s_{t+1}$
- 6:         **Value update (parameter update):**

$$w_{t+1} = w_t + \alpha_t \left[ r_{t+1} + \gamma \max_{a \in \mathcal{A}(s_{t+1})} \hat{q}(s_{t+1}, a, w_t) - \hat{q}(s_t, a_t, w_t) \right] \nabla_w \hat{q}(s_t, a_t, w_t)$$

- 7:         **Policy update:**

$$\pi_{t+1}(a \mid s_t) = \begin{cases} 1 - \frac{\varepsilon}{|\mathcal{A}(s)|} (|\mathcal{A}(s)| - 1) & \text{if } a = \arg \max_{a \in \mathcal{A}(s_t)} \hat{q}(s_t, a, w_{t+1}) \\ \frac{\varepsilon}{|\mathcal{A}(s)|} & \text{otherwise} \end{cases}$$

# Motivation for off-policy learning

- In X-ray Computed Tomography (CT), projections from many angles are acquired and used for 3D reconstruction.<sup>a</sup> To make CT suitable for in-line quality control, reducing the number of angles while maintaining reconstruction quality is necessary.
- However, it is expensive and even unethical to employ different policies on patients for the purpose of training an optimized policy.
- **off-policy learning**: learning from experiences generated by a different policy than the one it's currently following.




Fig. 1. The interaction between the environment and the agent during policy training

<sup>a</sup>Tianyuan Wang, Felix Lucka, and Tristan van Leeuwen.

"Sequential experimental design for X-ray CT using deep reinforcement learning". In: *arXiv preprint arXiv:2307.06343* (2023).

# Deep Q-learning

- Deep Q-learning aims to optimize the following objective function:

$$J = \mathbb{E} \left[ \left( R + \gamma \max_{a \in \mathcal{A}(S')} \hat{q}(S', a, w) - \hat{q}(S, A, w) \right)^2 \right]$$

- where  $(S, A, R, S')$  are random variables that denote a state, an action, the immediate reward, and the next state, respectively.
- Note that the parameter  $w$  appears in  $R + \gamma \max_{a \in \mathcal{A}(S')} \hat{q}(S', a, w)$ , this brings some computational challenge as it is nontrivial to calculate the gradient.
- One potential method is to fix  $w$  when computing the gradient and update  $w$  every  $C$  iterations.
- We introduce two networks: one is a main network representing  $\hat{q}(s, a, w)$  and the other is a target network  $\hat{q}(s, a, w_T)$ . The objective function in this case becomes

$$J = \mathbb{E} \left[ \left( R + \gamma \max_{a \in \mathcal{A}(S')} \hat{q}(S', a, w_T) - \hat{q}(S, A, w) \right)^2 \right],$$

where  $w_T$  is the target network's parameter.

# Deep Q-learning

- Now we can easily compute the gradient. When  $w_T$  is fixed, the gradient of  $J$  is

$$\nabla_w J = -\mathbb{E} \left[ \left( R + \gamma \max_{a \in \mathcal{A}(S')} \hat{q}(S', a, w_T) - \hat{q}(S, A, w) \right) \nabla_w \hat{q}(S, A, w) \right]$$

---

## Algorithm Deep Q-learning (off-policy version)

---

- 1: **Initialization:** A main network and a target network with the same initial parameter.
  - 2: **Goal:** Learn an optimal target network to approximate the optimal action values from the experience samples generated by a given behavior policy  $\pi_b$ .
  - 3: Store the experience samples generated by  $\pi_b$  in  $\mathcal{B} = \{(s, a, r, s')\}$
  - 4: **for** each iteration **do**
  - 5:     Uniformly draw a mini-batch of samples from  $\mathcal{B}$
  - 6:     For each sample  $(s, a, r, s')$ , Calculate the target value as  
       $y_T = r + \gamma \max_{a \in \mathcal{A}(s')} \hat{q}(s', a, w_T)$ , where  $w_T$  is the parameter of the target network
  - 7:     Update the main network to minimize  $(y_T - \hat{q}(s, a, w))^2$  using the mini-batch of samples
  - 8:     Set  $w_T = w$  every  $C$  iterations
  - 9: **end for**
-

# Content

- 1: *Basic setting in reinforcement learning*
- 2: *Intro to Q-learning*
- 3: **Case study: Q-learning in PM**

# Q Learning in PM

## Multi-decision Setting<sup>1</sup>

- Recall that the observed data is of the form  $\{(X_{t,i}, A_{t,i}, Y_{t,i})\}_{i=1}^n$ . At each decision point  $t = 1, \dots, T$ , assume that there is a finite set of all possible treatment options  $\mathcal{A}_t$  with elements  $A_t \in \mathcal{A}_t$ .
- Let  $Y_T$  be the proximal outcome measured after the treatment at stage  $T$ .
- Denotes  $H_t$  as the set of available patient history at time  $t$  such that
  - $H_1 = X_1$
  - $H_t = (H_{t-1}, A_{t-1}, Y_{t-1}, X_t)$
- A dynamic Treatment Regime is a sequence of functions  $\mathbf{d} = (d_1, \dots, d_T)$  such that  $d_t : \mathcal{H}_t \rightarrow \mathcal{A}_t$  for  $t = 1, \dots, T$ .
- An optimal treatment regime maximizes the expectation of some (prespecified) cumulative outcome measure  $Y = y(Y_1, \dots, Y_T)$ , e.g.,  
 $y(v_1, \dots, v_T) = \sum_{t=1}^T v_t$ , or  $y(v_1, \dots, v_T) = \max_t v_t$ , or  $y(v_1, \dots, v_T) = v_T$ .

---

<sup>1</sup>Phillip J Schulte et al. "Q-and A-learning methods for estimating optimal dynamic treatment regimes". In: *Statistical science: a review journal of the Institute of Mathematical Statistics* 29.4 (2014), p. 640.



# Q Learning in PM

## Optimal Treatment Regime<sup>3</sup>

- Under these assumptions, we can express the optimal regimes in terms of the observed data. We now define the following:

$$Q_T(h_T, a_T) = E(Y_T \mid H_T = h_T, A_T = a_T)$$

$$V_T(h_T, a_T) = \max_{a_T} Q_T(h_T, a_T)$$

- and for  $t = T - 1, \dots, 1$ ,

$$Q_t(h_t, a_t) = E(V_{t+1}(h_{t+1}, a_t) \mid H_t = h_t, A_t = a_t)$$

$$V_t(h_t, a_t) = \max_{a_t} Q_t(h_t, a_t)$$

- The optimal DTRs is:

$$d_t^{\text{opt}}(h_t) = \arg \max_{a_t} Q_t(h_t, a_t), \quad \text{for } t = 1, \dots, T \quad (2)$$

- Q-learning is an approximate dynamic programming<sup>2</sup> algorithm based on (2). This immediately suggests a regression-based estimator  $\hat{Q}_{t,n}(h_t, a_t, \xi_t)$  of  $Q_{t,n}(h_t, a_t)$  by regressing  $Y$  on  $H_t$  and  $A_t$ , where  $\xi_t$  is the parameters for estimating  $\hat{Q}_{t,n}$ .

---

<sup>2</sup>Richard Bellman. "Dynamic programming". In: *Science* 153.3731 (1966), pp. 34–37.

<sup>3</sup>Schulte et al., "Q-and A-learning methods for estimating optimal dynamic treatment regimes".

# Q Learning in PM

## Q-learning for Two Stages<sup>4</sup>

- Considering DTR with only two stages. We may fit linear models for

$$Q_1(h_1, a_1; \xi_1) = \mathcal{H}_1^T \beta_1 + a_1 (\mathcal{H}_1^T \psi_1)$$

$$Q_2(h_2, a_2; \xi_2) = \mathcal{H}_2^T \beta_2 + a_2 (\mathcal{H}_2^T \psi_2)$$

- where

$$\mathcal{H}_1 = (1, x_1^T)^T \quad \mathcal{H}_2 = (1, x_1^T, a_1, x_2^T)^T$$

$$\xi_t = (\beta_t^T, \psi_t^T)^T \quad t = 1, 2$$

- Here  $Q_2(h_2, a_2; \xi_2)$  is a model for  $E(Y | H_2 = h_2, A_2 = a_2)$ , a standard regression problem involving observable data, whereas  $Q_1(s_1, a_1; \xi_1)$  is a model for  $E(V_2(h_2, a_1) | H_1 = h_1, A_1 = a_1)$

- The corresponding V-functions are

$$V_2(h_2, a_2; \xi_2) = \max_{a_2 \in \{-1, 1\}} Q_2(h_2, a_2; \xi_2)$$

$$= \mathcal{H}_2^T \beta_2 + (\mathcal{H}_2^T \psi_2) \times \text{sign}(\mathcal{H}_2^T \psi_2), \text{ and}$$

$$V_1(s_1; \xi_1) = \max_{a \in \{-1, 1\}} Q_1(s_1, a; \xi_1)$$

$$= \mathcal{H}_1^T \beta_1 + (\mathcal{H}_1^T \psi_1) \text{sign}(\mathcal{H}_1^T \psi_1)$$

- We can see that

$$d_1^{\text{opt}}(h_1; \xi_1) = \text{sign}(\mathcal{H}_1^T \psi_1)$$

$$d_2^{\text{opt}}(h_2, a_1; \xi_2) = \text{sign}(\mathcal{H}_2^T \psi_2)$$

- We can see that we only need to estimate the regression coefficients  $\psi_1$  and  $\psi_2$ , which can be done via OLS and WLS, etc.

<sup>4</sup>Schulte et al., "Q-and A-learning methods for estimating optimal dynamic treatment regimes".

# Q-learning for Multi-stages

- Here we further explore to multi-stage scenario. Consider the model  $Q_t(h_t, a_t; \xi_t)$ , say, for  $t = T, T-1, \dots, 1$ , each depending on a finite-dimensional parameter  $\xi_t$
- The models may be linear or nonlinear in  $\xi_t$  and include main effects and interactions in the elements of  $h_t$  and  $a_t$ .
- Estimators  $\hat{\xi}_t$  can be obtained in a backward iterative fashion for  $t = T, T-1, \dots, 1$  by solving suitable estimating equation (e.g., ordinary (OLS) or weighted (WLS) least squares)

$$\sum_{i=1}^n \frac{\partial Q_t(H_{ti}, A_{ti}; \xi_t)}{\partial \xi_t} \Sigma_t^{-1}(H_{ti}, A_{ti}) \left\{ \hat{V}_{(t+1)i} - Q_t(H_{ti}, A_{ti}; \xi_t) \right\} = 0$$

- where  $\Sigma_t(h_t, a_t)$  is a variance covariance matrix and  $\hat{V}_{ti}$  is an estimation of  $V_{ti}$  by incorporating the  $\hat{\xi}_t$  across all  $i = 1, \dots, n$ . For  $t = T$ , letting  $\hat{V}_{(T+1)i} = Y_i$ .

# Q-learning for Multi-stages

- We can then optimize  $Q_t$  by substituting  $\hat{\xi}_t$  for  $\xi_t$  and maximizing the equation

$$d_t^{\text{opt}} \left( h_t, a_{t-1}; \hat{\xi}_T \right) = \arg \max_{a_t \in \Psi_t(h_t)} Q_t \left( h_t, a_{t-1}, \hat{\xi}_t \right)$$

where  $\Psi_t(h_t)$  represents the set of allowable treatments for patients at the  $t$ -th treatment stage given available patient history at time  $t$ .

- For each  $i$ , update

$$\hat{V}_{ti} = \max_{a_t \in \Psi_t(H_{ti})} Q_t \left( H_{ti}, A_{(t-1)i}, a_t; \hat{\xi}_t \right)$$

- Summarize the estimated optimal regime as  $\hat{d}_Q^{\text{opt}} = \left( \hat{d}_{Q,1}^{\text{opt}}, \dots, \hat{d}_{Q,T}^{\text{opt}} \right)$

$$\hat{d}_{Q,1}^{\text{opt}}(h_1) = d_1^{\text{opt}}(h_1; \hat{\xi}_1)$$

$$\hat{d}_{Q,t}^{\text{opt}}(h_t, a_{t-1}) = d_t^{\text{opt}}(h_t, a_{t-1}; \hat{\xi}_t)$$

Thank you!

You can contact us at:

[yangruan@unc.edu](mailto:yangruan@unc.edu)   [yuananyan@unc.edu](mailto:yuananyan@unc.edu)   [yajie@unc.edu](mailto:yajie@unc.edu)